

Backend Development Roadmap

Phase 1: Core Node.js Mastery (Week 1–2)

Key Concepts

- **Asynchronous Programming**
 - Callback functions
 - Promises
 - `async/await` patterns
- **Event Loop & Concurrency**
 - Call Stack, Task Queue, Event Queue
 - Event loop phases and timers
- **Streams**
 - Types: Readable, Writable, Duplex, Transform
 - Use cases: Large file processing, data piping
- **Error Handling**
 - Error-first callback style
 - `try/catch` with `async/await`
 - Custom error classes
- **Performance Optimization**
 - Non-blocking code practices
 - Avoiding memory leaks
 - Monitoring performance bottlenecks
- **Cluster Module**
 - Forking child processes
 - Load balancing across CPU cores

- **Core Node Modules**

- `fs` , `path` , `os` , `crypto` , `url` , `buffer`

- **Memory Management**

- Detecting and debugging memory leaks
- Tools: `clinic.js` , Chrome DevTools heap profiling

Phase 2: Express.js Advanced (Week 3–4)

Core Concepts

- **Express.js Fundamentals Refresher**

- How Express handles HTTP requests/responses
- Request-Response lifecycle

Middleware

- **Types of Middleware**

- Application-level
- Router-level
- Error-handling middleware
- Built-in vs. Custom middleware

- **Practical Use Cases**

- Logging, body parsing, authentication, error tracking

Security & Protection

- **CORS (Cross-Origin Resource Sharing)**

- Configuring CORS with `cors` package
- Pre-flight requests and headers

- **Helmet.js**

- Setting HTTP headers for security
- Preventing XSS, clickjacking, MIME sniffing

- **Rate Limiting & Throttling**

- Using `express-rate-limit` and `Redis` to prevent DDoS attacks
- **Session Management**
 - Sessions vs. JWT
 - Implementing Redis-backed sessions
 - Cookie settings and security (HttpOnly, Secure, SameSite)

Input Handling

- **Validation & Sanitization**
 - Using `express-validator`
 - Escaping user input
- **Custom Validators**
 - Writing schema-specific custom rules

API Architecture

- **API Versioning Techniques**
 - URI versioning: `/api/v1/`
 - Header-based versioning
 - **Organizing Route Handlers**
 - Modular routers
 - Error responses structure (standardizing errors)
-

Phase 3: MongoDB Deep Dive (Week 5–6)

NoSQL vs SQL

- Key differences: schema, joins, normalization, scalability
- When to choose MongoDB over relational databases

Schema Design with Mongoose

- **Defining Mongoose Models**
 - Data types, default values, required fields

- **Validation & Middleware**
 - Pre/post hooks
 - Async custom validations

Query Performance

- **Indexes**
 - Types: single field, compound, text, hashed
 - Performance considerations
- **Aggregation Pipeline**
 - `$match` , `$group` , `$lookup` , `$unwind`
 - Building reports and analytics

Advanced MongoDB Features

- **Replication**
 - Primary-secondary setup
 - Failover scenarios
- **Sharding**
 - Horizontal partitioning across clusters
 - Use cases and limitations
- **Transactions**
 - Multi-document ACID operations
 - Sessions and rollback scenarios

Data Integrity & Optimization

- Ensuring consistency in denormalized data
 - Query profiling using MongoDB Compass and `explain()`
 - Best practices in schema design and query building
-

Phase 4: RESTful API Architecture (Week 7–8)

API Design Patterns

- **RESTful Principles**
 - Statelessness, uniform interface, layered system
- **Resource Naming Conventions**
 - Nouns vs Verbs in URIs
 - Versioning and pagination

CRUD Implementation

- Create, Read, Update, Delete using Express and Mongoose
- Soft deletes, timestamps, and audit trails

API Documentation

- Using **Swagger / OpenAPI**
 - Auto-generate docs
 - Interactive API testing via Swagger UI

Request Validation

- Middleware for payload schema checking
- Centralized validation error handling

Error Handling Strategy

- Consistent error structure (code, message, stack)
- HTTP status codes (200, 400, 401, 404, 500)

Testing the APIs

- Unit Testing with **Jest**
- API Testing with **Supertest**
- Mocking database calls for isolated tests

Phase 5: Authentication & Authorization (Week 9–10)

Authentication Basics

- **What is Authentication vs Authorization**
 - Difference and importance
 - Real-world examples and user flows

JWT (JSON Web Tokens)

- Structure of JWT: Header, Payload, Signature
- Generating & signing tokens with `jsonwebtoken`
- Verifying tokens and protecting routes
- Expiry, refresh tokens, and token rotation

OAuth2 Integration

- Using third-party auth (Google, Facebook)
- Setting up OAuth2 flows with Passport.js
- Access tokens vs ID tokens
- Handling OAuth callbacks and state

Role-Based Access Control (RBAC)

- Defining roles (admin, user, guest)
- Role permissions in route guards
- Middleware for role checks

Advanced Security

- **2FA (Two-Factor Authentication)**
 - Generating and verifying OTPs
 - Sending OTPs via Email/SMS (using Twilio/Nodemailer)
- **Password Hashing with Bcrypt**
 - Storing hashed passwords securely
 - Salting and comparing passwords
- **Secure Password Reset Flows**

- Token-based password reset mechanism
- Emailing reset links securely

Session & Cookie Management

- Storing sessions in Redis
 - HTTP-only, secure cookies with session tokens
 - Expiry handling and logout mechanisms
-

Phase 6: Testing & Debugging (Week 11–12)

Types of Testing

- Unit Testing: Test individual logic components (functions, utils)
- Integration Testing: Test routes and controllers
- End-to-End Testing: Simulate real user flows (Optional)

Testing Tools

- **Jest**
 - Writing test cases with `describe` , `it` , `expect`
 - Mocking functions and modules
- **Supertest**
 - Testing REST API endpoints
 - Setting up HTTP request-response scenarios

Mocking & Isolation

- Jest mocks and spies
- Mocking Mongoose and third-party libraries
- Avoiding real DB/API calls in tests

Test-Driven Development (TDD)

- Writing tests before functionality
- Red-green-refactor workflow

- Test coverage metrics

Debugging Techniques

- Node Inspector / Chrome DevTools
- VS Code Debugger setup
- Debugging async/await and promises

Continuous Testing & Automation

- Setting up **GitHub Actions** or **Travis CI**
 - Auto-run tests on push/PR
 - Coverage reports using tools like Coveralls or Codecov
-

Phase 7: Performance, Deployment & CI/CD (Week 13–14)

Performance Optimization

- Identifying bottlenecks with `clinic.js`, Node Profiler
- Caching (Redis) strategies for expensive DB queries
- Async optimizations and load testing (e.g., Artillery)

Clustering & Load Balancing

- Using **PM2** for process management
- Scaling Node.js apps using clustering
- Sticky sessions and Redis session sharing

Docker & Containers

- Writing Dockerfiles for Node.js projects
- Using `.dockerignore`, `docker-compose.yml`
- Building, running, and networking containers

CI/CD Pipelines

- Using **GitHub Actions** or **Jenkins**

- Build-test-deploy cycles automation
- Linting and formatting automation with ESLint + Prettier

Deployment

- Platforms: **Render**, **Railway**, **Vercel**, **AWS EC2**, or **DigitalOcean**
- Managing `.env` files with environment variables
- Handling production errors/logs using tools like Sentry

WebSockets Integration

- Real-time data (chat, notifications)
- Using **Socket.IO** for WebSocket communication

Rate Limiting Revisited

- Protecting endpoints from high-traffic abuse
- IP-based, user-based limits with Redis

Phase 8: Containerization & Microservices (Week 15–16)

Docker Essentials

- **Docker Architecture:** Images, containers, registries, volumes, networks
- **Docker CLI:** `build`, `run`, `exec`, `logs`, `prune`, `rm`, etc.
- Creating Dockerfiles for Node.js projects
- Best practices: Alpine base image, multi-stage builds, non-root users
- Building a production-ready container for a Node + Mongo app

Docker Compose

- `docker-compose.yml` basics
- Running multi-container apps (Node.js + MongoDB + Redis)
- Environment variable management
- Volume mounts and service networks

Introduction to Microservices

- **Monolith vs Microservices:** Pros, cons, and tradeoffs
- Breaking down a monolith into microservices
- Designing independent services with bounded contexts
- Keeping microservices stateless

API Gateway

- Introduction to API Gateway pattern
- Tools: Nginx as reverse proxy or using **Kong**
- Centralized authentication and request routing

Service Discovery

- Why it matters in microservices
- Tools: **Consul**, **Eureka**, DNS-based discovery
- Dynamic registration and deregistration of services

Event-Driven Architecture

- Benefits of asynchronous communication
- **Message brokers** overview: Kafka vs RabbitMQ
- Event producers, consumers, and message queues

Kafka Basics

- What is Kafka and when to use it
- Kafka topics, partitions, producers, and consumers
- Kafka in Node.js using `kafkajs`
- Handling retries, dead-letter queues, and message durability

Inter-Service Communication

- RESTful APIs between services
 - Kafka or RabbitMQ for pub-sub messaging
 - gRPC: Setup, Protobufs, and performance benefits
-

Phase 9: GraphQL & Kafka Deep Dive (Week 17–18)

GraphQL Basics

- What is GraphQL and how it differs from REST
- **Schemas:** Defining types, queries, mutations
- Query nesting, arguments, and aliases
- Apollo Playground for testing queries

Apollo Server with Node.js

- Setting up Apollo Server
- Modular schema design using `makeExecutableSchema`
- Context API for injecting authentication/authorization
- Connecting GraphQL to MongoDB with resolvers

GraphQL vs REST

- When to use GraphQL (complex querying, flexibility)
- REST advantages (simplicity, caching)
- Performance considerations, rate limiting, tooling

GraphQL Subscriptions (Real-Time Updates)

- Using WebSocket or `graphql-ws`
- Subscribing to real-time events (e.g., chat messages, notifications)
- Publishing from backend services using PubSub

Kafka in Production

- **Kafka Setup:** Zookeeper, brokers, topics
- Consuming Kafka in Node.js using `kafkajs` or `node-rdkafka`
- Efficient batch processing of messages
- Kafka Connect for data sync (MongoDB to Kafka)

Event Sourcing

- Events as the single source of truth
- Storing events instead of state
- Rebuilding state from event logs
- Benefits: audit trail, replayable workflows

Kafka Streams (Optional Advanced)

- Real-time data processing with Kafka Streams API
- Filtering, joining, and transforming message streams

GraphQL Federation (Optional Advanced)

- Federating multiple GraphQL services
- Apollo Gateway setup
- Shared ownership of the schema across teams

Interview Preparation Topics (Week 19–20)

1. Data Structures & Algorithms (DSA)

Mastering DSA is essential not just for coding interviews but also to improve backend logic and optimization.

Core Concepts:

- **Arrays & Strings:** Sliding window, two-pointer techniques
- **Linked Lists:** Reversals, cycles, merging
- **Stacks & Queues:** Balanced parentheses, infix-postfix conversion
- **Trees & Binary Trees:** Inorder/preorder/postorder traversals, BST, AVL Trees
- **Graphs:** DFS, BFS, Dijkstra's, Union-Find, Topological Sort
- **Heaps:** Min/Max Heap, Priority Queue use cases
- **Recursion & Backtracking:** Combinatorics, permutations, subset generation
- **Dynamic Programming:** Memoization, Tabulation, 0/1 Knapsack, LIS

- **Searching & Sorting:** Binary search, Merge sort, Quick sort, Radix sort

Practice Platforms:

- [LeetCode](#)
 - [HackerRank](#)
 - [Codeforces](#)
 - [InterviewBit](#)
-

2. Design Patterns

Understanding reusable patterns enhances your backend software architecture.

- **Singleton Pattern:** For shared database connections or caching
- **Factory Pattern:** Abstract object creation for services or strategies
- **Observer Pattern:** Event-based systems, WebSockets
- **Strategy Pattern:** For switching between algorithms at runtime
- **Decorator Pattern:** Extend functionality of classes or endpoints
- **Adapter Pattern:** Integration of third-party tools or legacy systems

Use Refactoring.Guru for crisp visual explanations.

3. SOLID Principles (Object-Oriented Design)

Helps keep backend code **clean, maintainable, and scalable**:

- **S – Single Responsibility:** One class, one job
 - **O – Open/Closed:** Extend, don't modify
 - **L – Liskov Substitution:** Subtypes replace parent types
 - **I – Interface Segregation:** Fine-grained interfaces
 - **D – Dependency Inversion:** Rely on abstractions, not concrete classes
-

4. Behavioral Interview Prep

Use the **STAR Method** to structure your answers:

- **Situation:** Set the context

- **Task:** What was your goal?
- **Action:** What actions did *you* take?
- **Result:** What was the outcome?

✅ **Sample Questions:**

- Tell me about a time you failed at something.
 - How do you handle team conflict?
 - Describe a challenging backend project.
 - Explain a time when you optimized a slow API.
-

5. Scalability & High Availability

Backend engineers must understand how to scale applications efficiently.

- **Horizontal Scaling:** Load balancing with Nginx, HAProxy
 - **Vertical Scaling:** Increasing server resources
 - **Replication:** MongoDB replicas for high availability
 - **Caching:** Redis, CDN usage
 - **Auto-scaling:** Cloud-native tools (AWS EC2 Auto Scaling)
-

6. High-Traffic API Management

Design APIs that handle **millions of requests** efficiently.

- **Rate Limiting:** Prevent abuse (e.g., 100 requests/min)
 - **Throttling:** Control over burst traffic
 - **Load Testing Tools:** Apache JMeter, Artillery, k6
 - **Circuit Breakers:** Fallbacks during service failures (e.g., Netflix Hystrix)
-

7. System Design (Big Picture Thinking)

Prepare to build **scalable backend systems** from scratch.

Key Topics:

- **Designing Scalable APIs:** Rate limits, retries, pagination

- **Design a URL Shortener:** Hashing, caching, database sharding
- **Design a Chat App:** WebSockets, message queues, delivery acknowledgment
- **Design an E-Commerce Backend:** Cart system, payment flows, order tracking
- **Design a Notification System:** Kafka + Webhooks or Push Services

Must-read Resources:

- [System Design Primer \(GitHub\)](#)
 - [ByteByteGo](#)
-



8. Microservices Design Patterns

- **Service Registry & Discovery:** Consul, Eureka
- **Circuit Breaker Pattern:** Netflix Hystrix, Resilience4j
- **Event Sourcing:** Kafka + immutable event logs
- **SAGA Pattern:** Distributed transaction management
- **Bulkhead Pattern:** Isolate failures across services
- **Centralized Logging & Monitoring:** ELK stack, Grafana + Prometheus